

Fiche TP 01:Linear Regression

Nom: Boudelal Abderahmane Amine & Nahoui Seif

Link colab to code exo1:

[TPO1_EXO 1.ipynb - Colab](#)

Exercice 1 :

Reponse1 :


```
: import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import SGDRegressor
```

Reponse 1: generate 2 vectors x and y :

generate two random vectors x and y wuth shape of (100,1)(100 rows and one columns)

```
: x = 2*np.random.rand(100,1)
y = 4+3*x+np.random.randn(100,1)*0.5
```

Reponse 2:



Reponse 3 :

Reponse 3:

SGDRegressor create a Stochastic Gradient Descent (SGD) with 100 iterations and small learning rate 0.0001 to avoid overfitting : model.fit to training the model

```
: model = SGDRegressor(max_iter=100,eta0=0.0001)
model.fit(x,y)
```

Reponse 4 :

reponse 4: Calculate Accuracy :

calculate the coefficient of determination using model.score or r2_score

```
from sklearn.metrics import r2_score
r2 = model.score(x,y)
print(f"the accuracy is :{r2}")
```

the accuracy is :-9.314908207188688

Reponse 5:


reponse 5 : interpret the resault

The negative R-squared value of -9.314908207188688 indicates a very poor model fit. This means the model's predictions are significantly worse than simply predicting the mean of the target variable. Causes: Poor Model Data issues Hyperparamater tuning we don't optimize our parameter in model

Reponse 6 :

The graph indicate an underfitting

```
[30]: plt.scatter(x,y)
plt.plot(x,model.predict(x),c='red')
plt.show()
```




Reponse 7 :

```
model = SGDRegressor(max_iter=1000,eta=0.001)
```

Reponse 8 :

```
#question 8
x_new = 2*np.random.rand(100,1)
model.fit(x,y.ravel())
score = model.score(x,y)
print('the score is',score)
#le plot
plt.scatter(x,y)
plt.plot(x_new, model.predict(x_new), color= 'red')
plt.show()
```

the score is 0.8448108524884594



✓ 0s — completed at 0:20.5

```
[12] r2 = r2_score(y,model.predict(x))  
     print(f"the accuracy is :{r2}")
```

⇒ the accuracy is :0.8448108524884594


Reponse 9

```

16] from io import StringIO
import sys
old_stdout = sys.stdout
sys.stdout = mystdout = StringIO()
clf = SGDRegressor(max_iter= 1000, alpha= 0.0001, verbose=1)
clf.fit(x,y.ravel())
sys.stdout = old_stdout
loss_history = mystdout.getvalue()
loss_list = []
for line in loss_history.split('\n'):
    if(len(line.split("loss: ")) == 1):
        continue
    loss_list.append(float(line.split("loss: ")[-1]))
plt.figure()
plt.plot(np.arange(len(loss_list)), loss_list)
plt.xlabel("Time in epochs")
plt.ylabel("Loss")

```

Text(0, 0.5, 'Loss')



Exercise 2 :

The link to the colab: [TP01_Exo02.ipynb - Colab](#)

Reponse 1 :

```
reponse 1: import libraries and dataset

: import numpy as np
import matplotlib.pyplot as plt
from sklearn import datasets, linear_model, model_selection

: x,y = datasets.load_diabetes(return_X_y=True)
```

REPOSNE 2

```
Reponse 2:

x.shape / print the data shape x[0]: First element

: print(x.shape)
print(y.shape)

(442, 10)
(442,)

: x[0]

: array([ 0.03807591,  0.05068012,  0.06169621,  0.02187239, -0.0442235 ,
        -0.03482076, -0.04340085, -0.00259226,  0.01990749, -0.01764613])
```

Réponse 3 :

```
Reponse 3 : ¶

x = x[:, 2] # Access the third column
x = x[:, np.newaxis]
```

Reponse4 :

```
Reponse 4:

split the data into train and test

x_train , x_test , y_train , y_test = model_selection.train_test_split(x,y,test_size=0.25,random_state=5)
```

Reponse 5 :

Reponse5:

create a linear regression Model

```
] : model = linear_model.LinearRegression()  
    model.fit(x_train,y_train)
```

```
] : LinearRegression ⓘ ?  
    LinearRegression()
```

Réponses 6 :

Reponses 6:

predict the data

```
predicted_y = model.predict(x_test)
```

Réponses 7 :

Reponses 7:

Accuracy score

```
: print(f"predicted y2 = {predicted_y[1]}")  
  print(f"actual y2 = {y_test[1]}")
```


```
predicted y2 = 222.3340825724795  
actual y2 = 332.0
```

```
: from sklearn.metrics import r2_score  
  accuracy = r2_score(y_test, predicted_y)  
  print(f"Accuracy {accuracy}")
```

```
Accuracy 0.33275881449507416
```

Réponses 8 :

```
plt.scatter(x_test,y_test)  
plt.plot(x_test,predicted_y, color= 'red')  
plt.show()
```



Réponses 9 :

Reponse 9 :

repeat stpes from 2 to 7 obtain most predictive feature is feature 9 with an R-squared score of 0.4215743297548834

```
] : # Load the diabetes dataset
X, y = datasets.load_diabetes(return_X_y=True)

# Iterate through each feature (except the first two)
best_feature = None
best_accuracy = 0
for i in range(2, X.shape[1]):
    X_feature = X[:, i].reshape(-1, 1) # Extract the current feature
    X_train, X_test, y_train, y_test = train_test_split(X_feature, y, test_size=0.25, random_state=42)
    model = LinearRegression()
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    accuracy = r2_score(y_test, y_pred)
    print(f"Accuracy for feature {i+1}: {accuracy}")
    if accuracy > best_accuracy:
        best_accuracy = accuracy
        best_feature = i+1

print(f"The most predictive feature is feature {best_feature} with an R-squared score of {best_accuracy}")
```

```
Accuracy for feature 3: 0.3172099449537781
Accuracy for feature 4: 0.18471383018307075
Accuracy for feature 5: 0.06626553266368607
Accuracy for feature 6: 0.04163111831069033
Accuracy for feature 7: 0.1377356557444941
Accuracy for feature 8: 0.19257590747667053
Accuracy for feature 9: 0.4215743297548834
Accuracy for feature 10: 0.16593419294983613
The most predictive feature is feature 9 with an R-squared score of 0.4215743297548834
```